

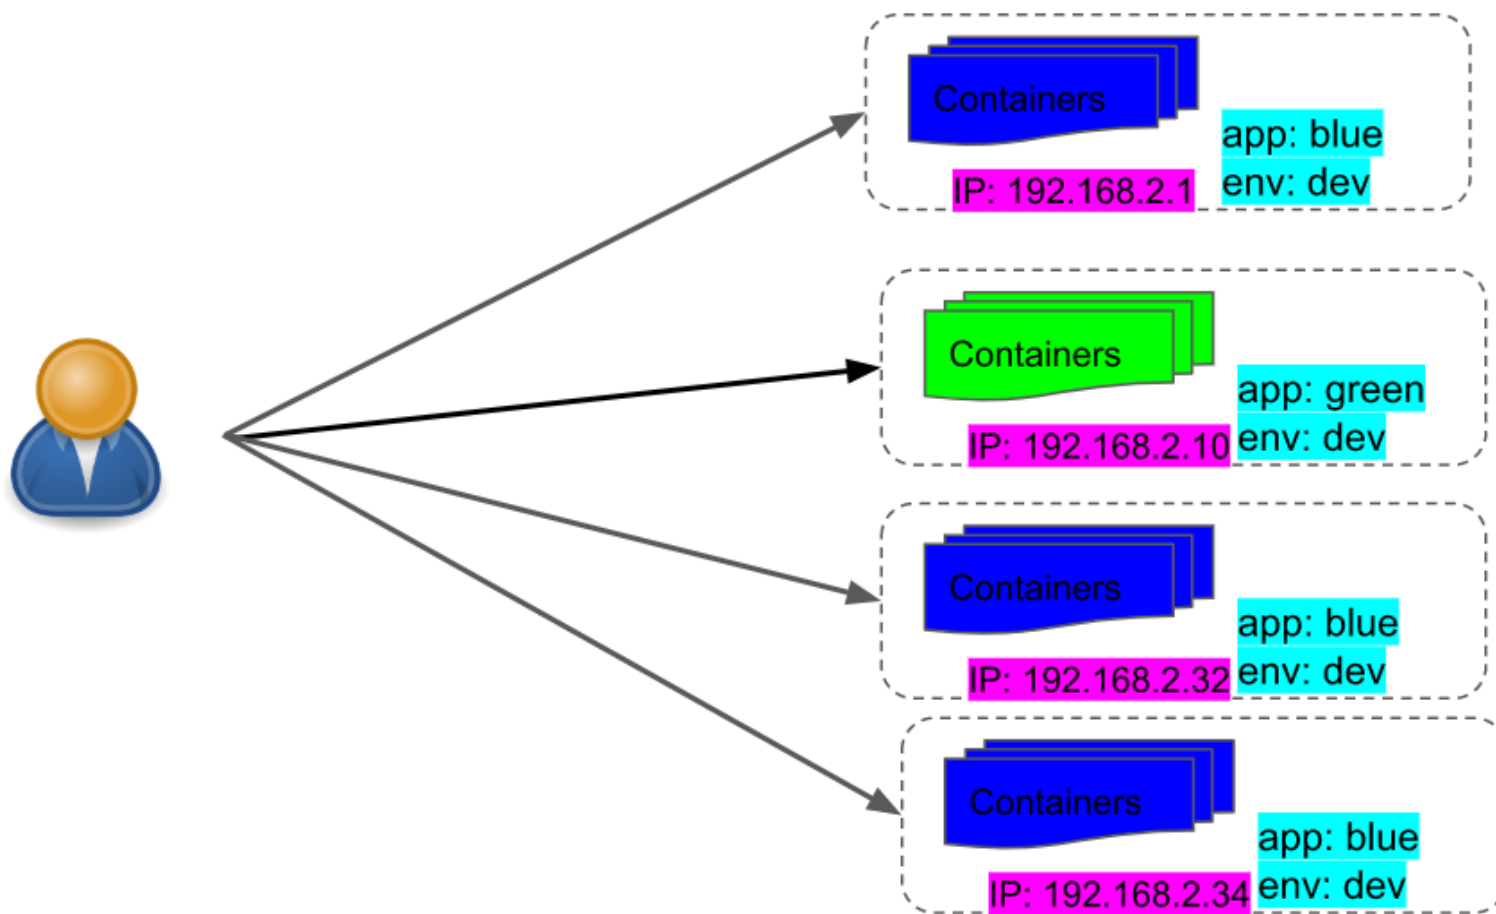
Kubernetes Services

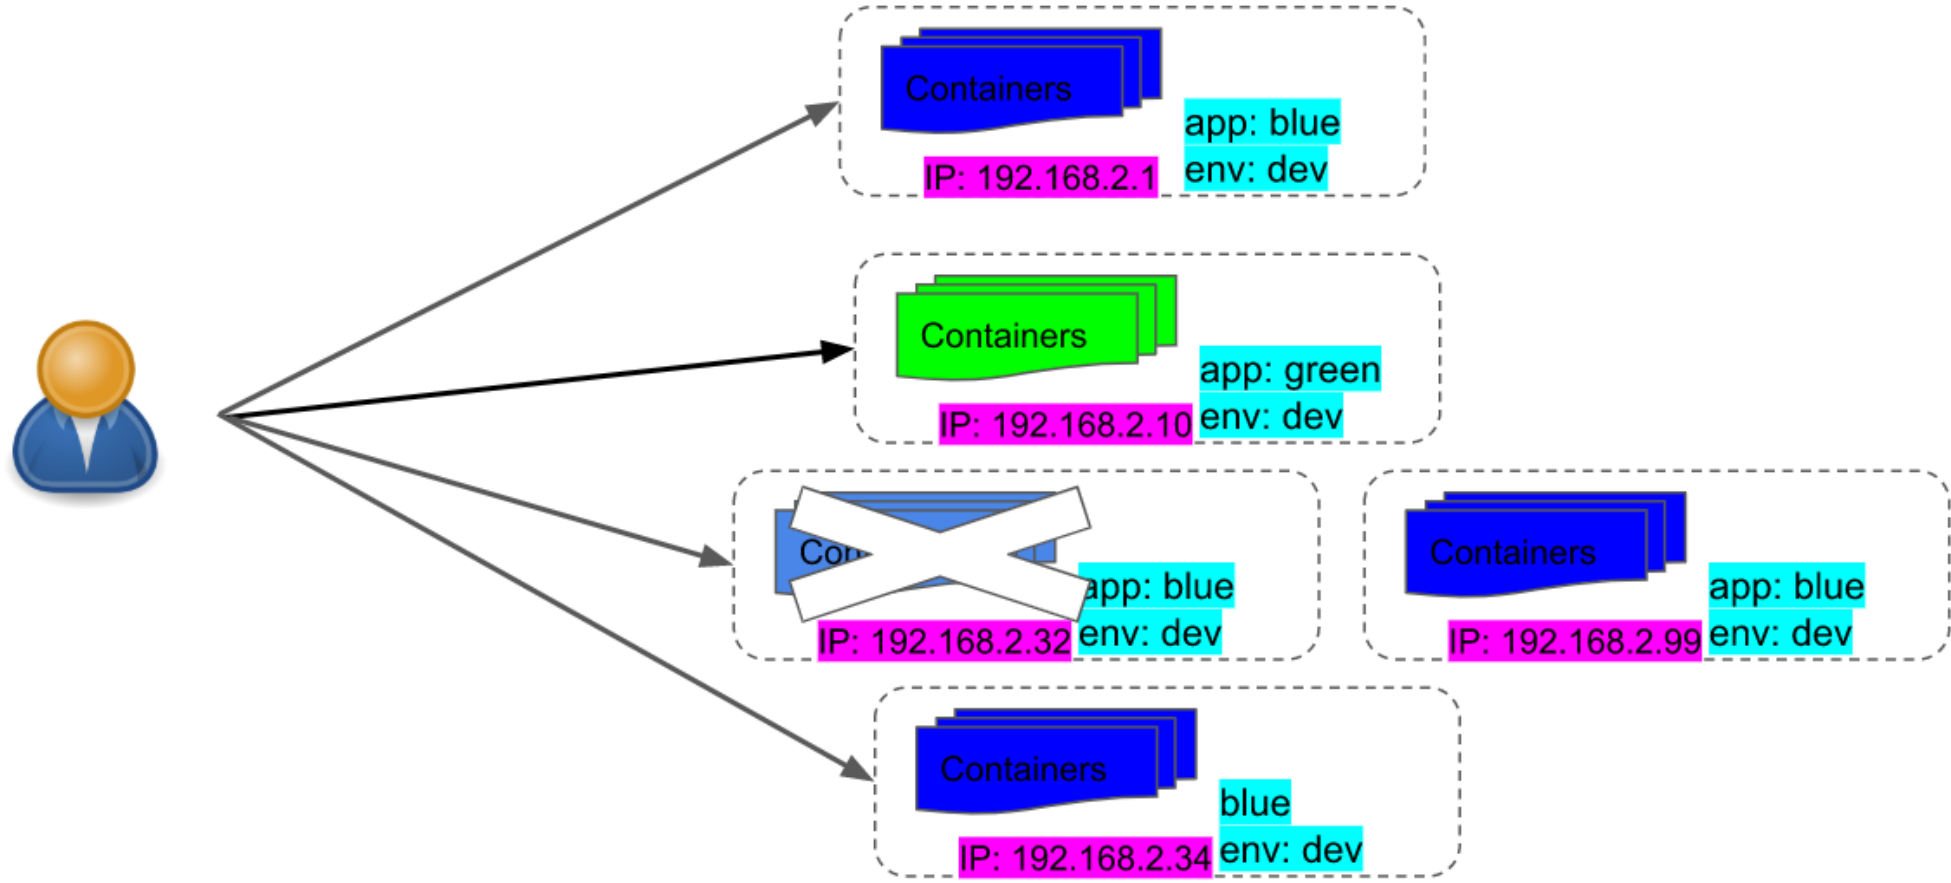
What is a Service?

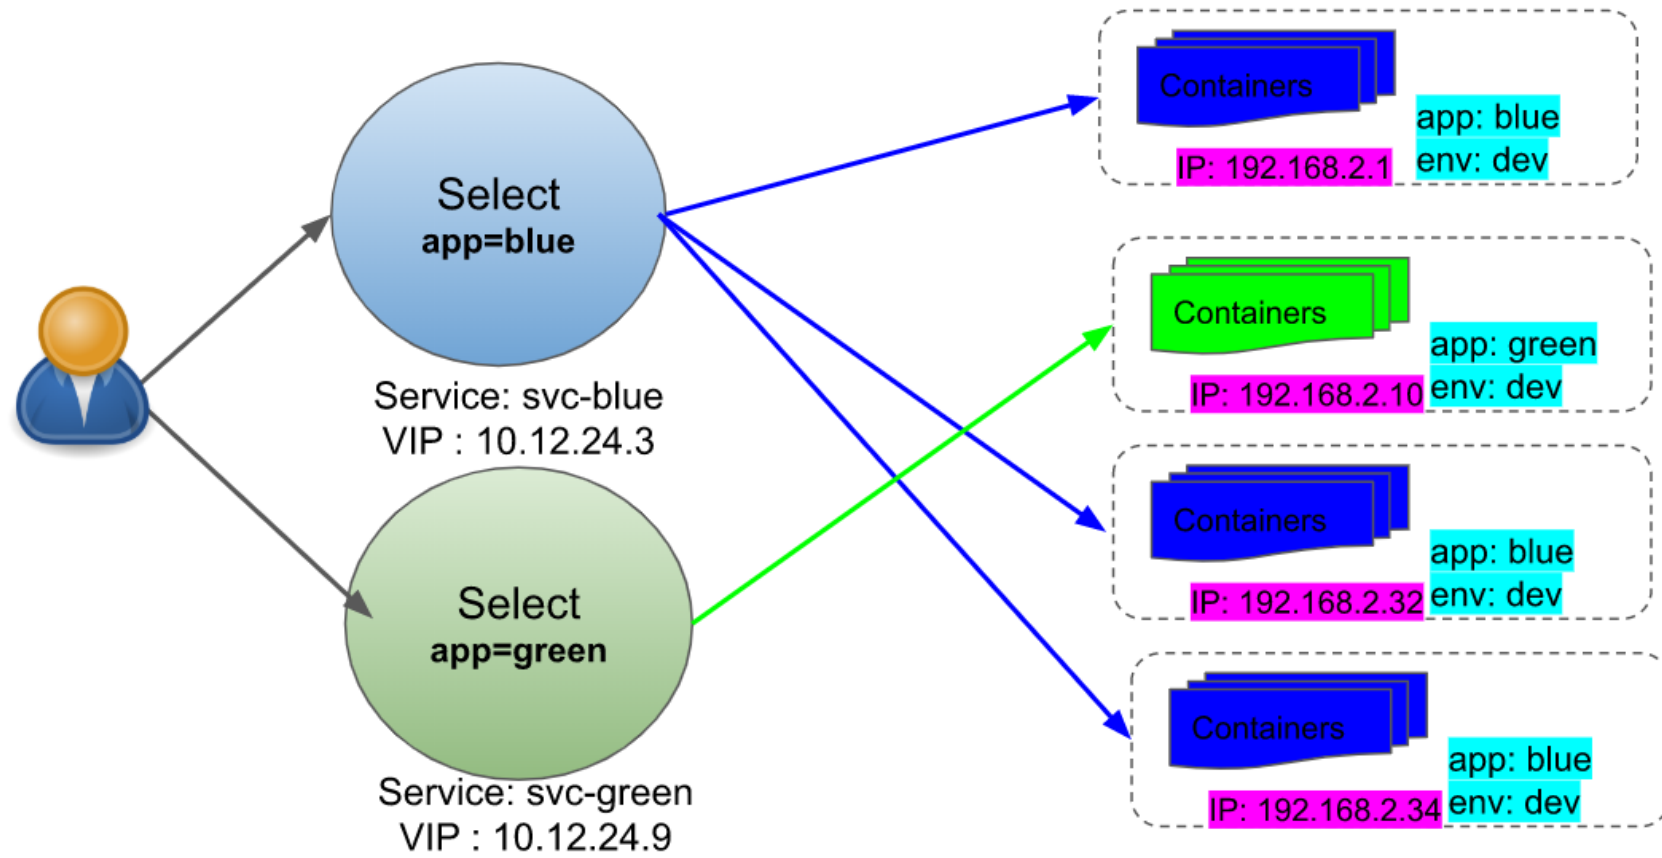
- ▶ A service can be defined as a logical set of pods.
- ▶ It can be defined as an abstraction on the top of the pod which provides a single IP address and DNS name by which pods can be accessed.
- ▶ With Service, it is very easy to manage load balancing configuration. It helps pods to scale very easily.
- ▶ A service is a REST object in Kubernetes whose definition can be posted to Kubernetes apiServer on the Kubernetes master t

Need of Kubernetes Services

- ▶ Kubernetes Pods details constantly changes . They may die and be restarted anywhere in the cluster at any time.
- ▶ When restarted, a Pod may be assigned a different IP address. This leads to a problem: if some set of Pods (let's call them backends) provides functionality to other Pods (let's call them frontends) inside the Kubernetes cluster, how do those frontends find out and keep track of which backends are in that set? Here, Services come to the rescue.
- ▶ A Kubernetes Service create a logical connection to the set of Pods and a policy by which to access them. The set of Pods targeted by a Service is usually determined by a Label Selector.

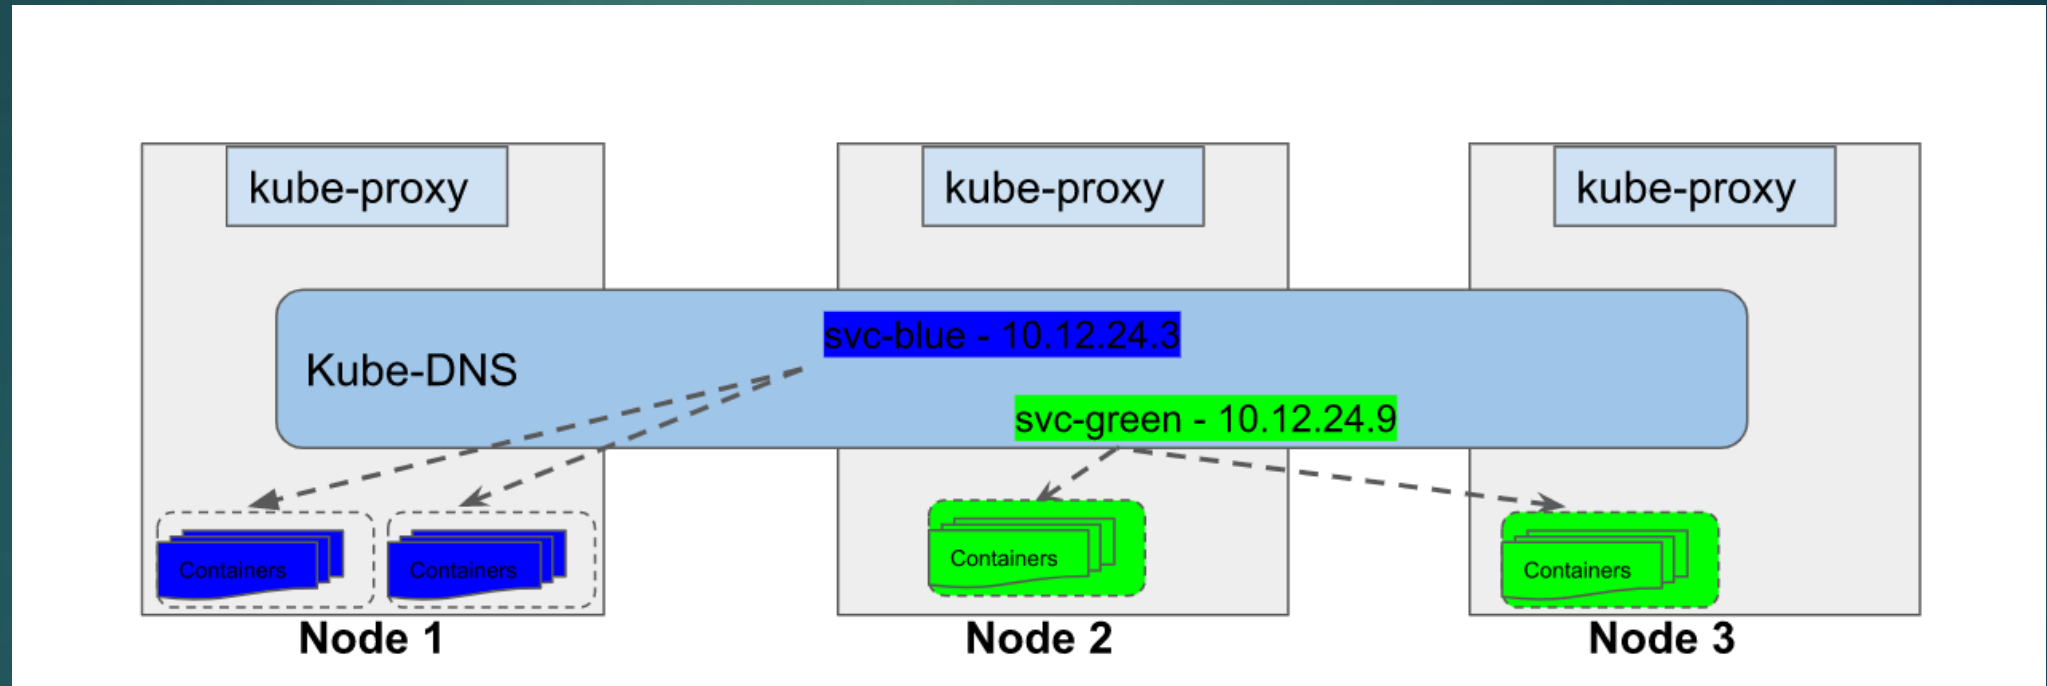




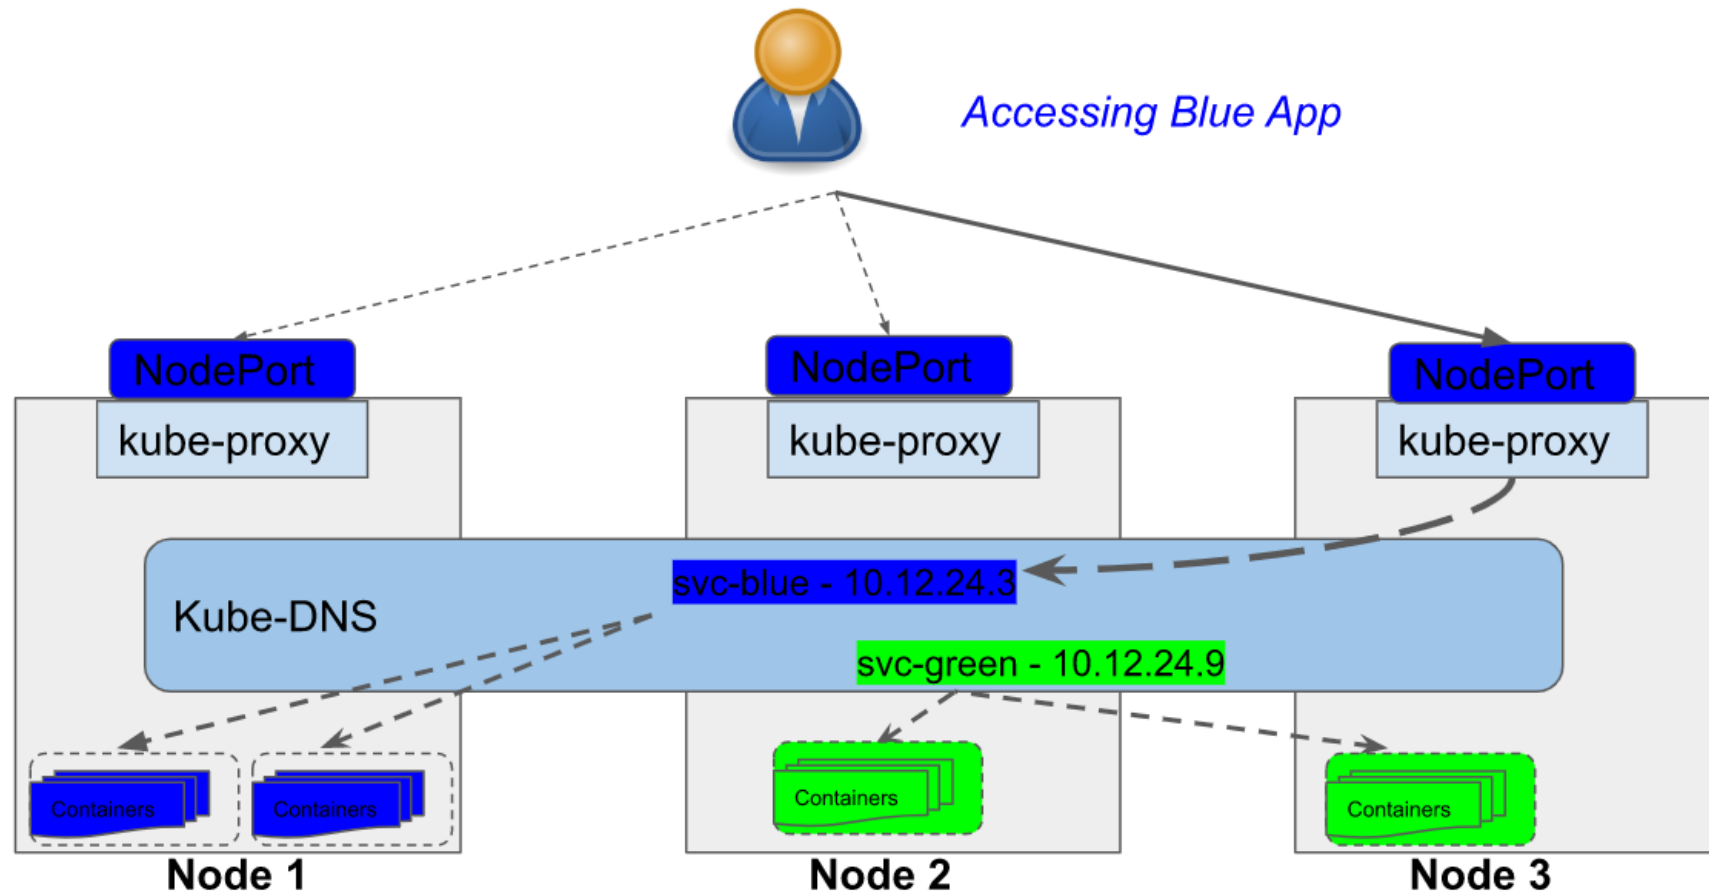


Cluster-Internal Ip (Default)

- ▶ This is the default service type. Choosing this value means that you want this service to be reachable only from inside of the cluster.



Service - NodePort



Service – Load Balancer

