

AWS Glacier

What is AWS Glacier

2

- ❑ Amazon Glacier is an extremely low-cost storage service that provides durable storage with security features for data archiving and backup.
- ❑ With Amazon Glacier, customers can store their data cost effectively for months, years, or even decades.
- ❑ You can upload a single file as an archive or aggregate multiple files into a TAR or ZIP file and upload as one archive.
- ❑ A single archive can be as large as 40 terabytes.
- ❑ Amazon Glacier is a REST-based web service.

AWS Glacier: Features

3

Low Cost

Secure

Durable


Simple

Flexible

Integrated

AWS Glacier: Data Model


4



AWS Glacier: Vault

5

- ❑ Vault is a container for storing archives.
- ❑ Within a region, an account must use unique vault names.
- ❑ An AWS account can create same-named vaults in different regions
- ❑ Each vault resource has a unique address. The general form is:
 - ❑ `https://<region-specific endpoint>/<account-id>/vaults/<vaultname>`
- ❑ Example, create a vault (examplevault) in the US West region.
 - ❑ `https://glacier.us-west-2.amazonaws.com/111122223333/vaults/examplevault`
 - glacier.us-west-2.amazonaws.com identifies the region.
 - 111122223333 is the AWS account ID that owns the vault.
 - vaults refers to the collection of vaults owned by the AWS account.
 - examplevault identifies a specific vault in the vaults collection.



© 2018 CHAITANYA GAJULA - ALL RIGHTS RESERVED

AWS Glacier: Archive

6

- ❑ Archive can be any data such as a photo, video, or document
- ❑ Archive is a base unit of storage in Amazon Glacier.
- ❑ Each archive has a unique ID.
- ❑ You can store an unlimited number of archives in a vault.
- ❑ Amazon Glacier assigns the archive an ID, which is unique in the AWS region in which it is stored.
 - ❑ `https://<region-specific endpoint>/<account-id>/vaults/`
 - ❑ `<vault-name>/archives/<archive-id>`
- ❑ You cannot update an existing archive.



© 2018 CHAITANYA GAJULA - ALL RIGHTS RESERVED

AWS Glacier: Archive Retrieval Options

Expedited

Allow you to quickly access your data
Data are typically made available within 1–5 minutes.

Types

- ✓ On-Demand
- ✓ Provisioned

Standard

Allow you to access your data within several hours
(Approx. 3-5 hours)

This is default option

Bulk

Lowest-cost retrieval option. Retrieve large amounts, of data inexpensively in a day.
Bulk retrievals complete within 5–12 hours.

AWS Glacier: Retrieval Policies

8

Free Tier

Max Retrieval Rate

No Retrieval Limit

AWS Glacier: Job

6

- ❑ To retrieve an archive, first initiate a job, and then download the job output after AWS Glacier completes the job.
- ❑ With AWS Glacier, your data retrieval requests are queued and most jobs take about four hours to complete.
- ❑ A vault can have multiple jobs in progress at any point in time.
- ❑ When you send a job request (initiate a job), AWS Glacier returns to you a job ID to track the job.
- ❑ For each job, AWS Glacier maintains information such as job type, description, creation date, completion date, and job status.
- ❑ You can obtain information about a specific job or obtain a list of all your jobs associated with a vault.
- ❑ Each job is uniquely identified in the form:
 - ❑ `https://<region-specific endpoint>/<account-id>/vaults/<vault-name>/jobs/<job-id>`

AWS Glacier: Notification Configuration

10

- ❑ Because jobs take time to complete, AWS Glacier supports a notification mechanism to notify you when a job is complete.
- ❑ You can configure a vault to send notification to an Amazon Simple Notification Service (Amazon SNS) topic when jobs complete.

AWS Glacier: Retrieval Options

11

❑ Glacier Pricing

- \$0.004 per GB / month

❑ Retrieval Pricing

Description	Expedited	Standard	Bulk
Data Access Time	1 - 5 minutes	3 - 5 hours	5 - 12 hours
Data Retrievals	\$0.03 per GB	\$0.01 per GB	\$0.0025 per GB
Retrieval Requests	On-Demand: \$0.01 per request Provisioned: \$100 per Provisioned Capacity Unit	\$0.050 per 1,000 requests	\$0.025 per 1,000 requests

AWS Glacier: Operation

12

Create Vault

Delete Vault

Retrieving Vault
Metadata

Downloading a
Vault Inventory

Uploading an
Archive


Downloading an
Archive

Deleting an
Archive

Updating an
Archive

AWS Glacier: Multipart Upload

13



AWS Glacier: Limitations

14

Description	Limitation
You can create up to 1,000 vaults per region	1000

THANK YOU